

Wei Bai

11431 Willows Rd NE #200,
Redmond, Washington, 98052 USA
✉ baiwei0427@gmail.com
📄 baiwei0427.github.io

Research Interests

I am broadly interested in computer networking, with special focuses on data center networking, congestion control, and high performance networked systems.

Education

- 2013.9-2017.6 **The Hong Kong University of Science and Technology (HKUST)**, Hong Kong SAR, China
Ph.D. Computer Science and Engineering
Advisor: Prof. Kai Chen
Thesis: "Congestion Control Mechanisms for Data Center Networks"
- 2009.9-2013.7 **Shanghai Jiao Tong University (SJTU)**, Shanghai, China
B.E. Information Security

Work Experience

- NVIDIA**, Redmond, USA
- 2024.1-Present Principal Software Research Architect, Networking Software and System Architecture Group
- Microsoft Research Redmond**, Redmond, USA
- 2021.7-2024.1 Senior Researcher, Networking Research Group
- 2020.9-2021.6 Senior Research Software Development Engineer, Networking Research Group
- 2019.8-2020.8 Research Software Development Engineer 2, Mobility and Networking Research Group
- Microsoft Research Asia**, Beijing, China
- 2017.7-2019.8 Researcher (formerly Associate Researcher 2), Networking Research Group

Selected Projects

- 2019.2-2024.1 **RoCE at Microsoft Azure**: Collaborated with Azure Networking, Storage, and Compute teams to design, develop, and operate high-performance and reliable network infrastructure for the world's largest remote direct memory access (RDMA) deployments, handling over 70% of Azure traffic. Key technology utilized was RDMA over Converged Ethernet (RoCE). Responsibilities included the qualification, testing, and optimization of all RoCE NICs and switches within the Azure infrastructure. Some operational experiences were shared in the paper "Empowering Azure Storage with RDMA" presented at NSDI 2023.

2018.9-Present **SONiC**: I have been contributing to Software for Open Networking in the Cloud (SONiC), an open source cross-platform switch operating system based on Linux. My contributions include support for disaggregated chassis, Quality of Service (QoS) and RDMA features, and development of a test system based on programmable hardware.

Awards and Honors

- 2021 Microsoft Special Stock Award
- 2016, 2017 HKUST Research Travel Grant
- 2015, 2016 USENIX NSDI Student Grant
- 2015 **Microsoft Research Asia Fellowship**
- 2013-2017 HKUST Postgraduate Scholarship
- 2013 SJTU Outstanding Bachelor Thesis Award (Top 1%)
- 2012 First Prize, Software Group, Information Security Contest, National Undergraduate Electronic Design Contest

Publications

Conference Publications

- [C30] Jiaqi Lou*, Xinhao Kong*, Jinghan Huang, **Wei Bai**, Nam Sung Kim, Danyang Zhuo, "Hardware-assisted RDMA Performance Isolation for Public Clouds", in *Proceedings of the 21st USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, Santa Clara, California, April 16–18, 2024.
- [C29] Vamsi Addanki, **Wei Bai**, Stefan Schmid, Maria Apostolaki, "Reverie: Low Pass Filter-Based Switch Buffer Sharing for Datacenters with RDMA and TCP Traffic", in *Proceedings of the 21st USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, Santa Clara, California, April 16–18, 2024.
- [C28] Hao Wang, Han Tian, Jingrong Chen, Xinchun Wan, Jiacheng Xia, Gaoxiong Zeng, **Wei Bai**, Junchen Jiang, Yong Wang, Kai Chen, "Towards Domain-Specific Network Transport for Distributed DNN Training", in *Proceedings of the 21st USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, Santa Clara, California, April 16–18, 2024.
- [C27] Zhuolong Yu, Bowen Su, **Wei Bai**, Shachar Raindel, Vladimir Braverman, Xin Jin, "Understanding the Micro-Behaviors of Hardware Offloaded Network Stacks with Lumina", in *Proceedings of the ACM SIGCOMM 2023 Conference (SIGCOMM)*, New York City, New York, September 10-14, 2023.
- [C26] Hwijoon Lim, Jaehong Kim, Inho Cho, Keon Jang, **Wei Bai**, Dongsu Han, "Flex-Pass: A Case for Flexible Credit-based Transport for Datacenter Networks", in *Proceedings of the 18th European Conference on Computer Systems (EuroSys)*, Rome, Italy, May 8-12 2023.

- [C25] **Wei Bai**, Shanim Sainul Abdeen, Ankit Agrawal, Krishan Kumar Attre, Paramvir Bahl, Ameya Bhagat, Gowri Bhaskara, Tanya Brokhman, Lei Cao, Ahmad Cheema, Rebecca Chow, Jeff Cohen, Mahmoud Elhaddad, Vivek Ette, Igal Figlin, Daniel Firestone, Mathew George, Ilya German, Lakhmeet Ghai, Eric Green, Albert Greenberg, Manish Gupta, Randy Haagens, Matthew Hendel, Ridwan Howlader, Neetha John, Julia Johnstone, Tom Jolly, Greg Kramer, David Kruse, Ankit Kumar, Erica Lan, Ivan Lee, Avi Levy, Marina Lipshteyn, Xin Liu, Chen Liu, Guohan Lu, Yuemin Lu, Xiakun Lu, Vadim Makhervaks, Ulad Malashanka, David A. Maltz, Ilias Marinos, Rohan Mehta, Sharda Murthi, Anup Namdhari, Aaron Ogus, Jitendra Padhye, Madhav Pandya, Douglas Phillips, Adrian Power, Suraj Puri, Shachar Raindel, Jordan Rhee, Anthony Russo, Maneesh Sah, Ali Sheriff, Chris Sparacino, Ashutosh Srivastava, Weixiang Sun, Nick Swanson, Fuhou Tian, Lukasz Tomczyk, Vamsi Vadlamuri, Alec Wolman, Ying Xie, Joyce Yom, Lihua Yuan, Yanzhao Zhang, Brian Zill, "Empowering Azure Storage with RDMA", in *Proceedings of the 20th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, Boston, Massachusetts, April 17-19, 2023 (Operational Systems Track).
- [C24] Xinhao Kong, Jingrong Chen, **Wei Bai**, Yechen Xu, Mahmoud Elhaddad, Shachar Raindel, Jitendra Padhye, Alvin R. Lebeck, Danyang Zhuo, "Understanding RDMA Microarchitecture Resources for Performance Isolation", in *Proceedings of the 20th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, Boston, Massachusetts, April 17-19, 2023.
- [C23] Bojie Li, Gefei Zuo, **Wei Bai**, Lintao Zhang, "1Pipe: Scalable Total Order Communication in Data Center Networks", in *Proceedings of the ACM SIGCOMM 2021 Conference (SIGCOMM)*, Virtual Conference, August 23-27, 2021.
- [C22] Hwijoon Lim, **Wei Bai**, Yibo Zhu, Youngmok Jung, Dongsu Han, "Towards Timeout-less Transport in Commodity Datacenter Networks", in *Proceedings of the 16th European Conference on Computer Systems (EuroSys)*, Virtual Conference, April 26-28, 2021.
- [C21] Shuihai Hu, **Wei Bai**, Gaoxiong Zeng, Zilong Wang, Baochen Qiao, Kai Chen, Kun Tan, Yi Wang, "Aeolus: A Building Block for Proactive Transport in Datacenters", in *Proceedings of the ACM SIGCOMM 2020 Conference (SIGCOMM)*, Virtual Conference, August 10-14, 2020.
- [C20] Qun Huang, Haifeng Sun, Patrick P. C. Lee, **Wei Bai**, Feng Zhu, Yungang Bao, "OmniMon: Re-architecting Network Telemetry with Resource Efficiency and Full Accuracy", in *Proceedings of the ACM SIGCOMM 2020 Conference (SIGCOMM)*, Virtual Conference, August 10-14, 2020.
- [C19] **Wei Bai**, Shuihai Hu, Kai Chen, Kun Tan, Yongqiang Xiong, "One More Config is Enough: Saving (DC)TCP for High-speed Extremely Shallow-buffered Datacenters", in *Proceedings of the 39th IEEE International Conference on Computer Communications (INFOCOM)*, Virtual Conference, July 6-9, 2020.
- [C18] Junxue Zhang, **Wei Bai**, Kai Chen, "Enabling ECN for Datacenter Networks with RTT Variations", in *Proceedings of the 15th International Conference on emerging Networking EXperiments and Technologies (CoNEXT)*, Orlando, Florida, December 9-12, 2019.

- [C17] Gaoxiong Zeng, **Wei Bai**, Ge Chen, Kai Chen, Dongsu Han, Yibo Zhu, Lei Cui, "Congestion Control for Cross-Datacenter Networks", in *Proceedings of the 27th IEEE International Conference on Network Protocols (ICNP)*, Chicago, Illinois, October 7-10, 2019.
- [C16] Xiaodong Yi, Junjie Wang, Jingpu Duan, **Wei Bai**, Chuan Wu, Yongqiang Xiong, Dongsu Han, "FlowShader: a Generalized Framework for GPU-accelerated VNF Flow Processing", in *Proceedings of the 27th IEEE International Conference on Network Protocols (ICNP)*, Chicago, Illinois, October 7-10, 2019.
- [C15] Bojie Li, Tianyi Cui, Zibo Wang, **Wei Bai**, Lintao Zhang, "SocksDirect: Datacenter Sockets can be Fast and Compatible" in *Proceedings of the ACM SIGCOMM 2019 Conference (SIGCOMM)*, Beijing, China, August 19-24, 2019.
- [C14] Yang Cheng, Dan Li, Zhiyuan Guo, Binyao Jiang, Jiaxin Lin, Xi Fan, Jinkun Geng, Xinyi Yu, **Wei Bai**, Lei Qu, Ran Shu, Peng Cheng, Yongqiang Xiong, Jianping Wu, "DLBooster: Boosting End-to-End Deep Learning Workflows with Offloading Data Preprocessing Pipelines" in *Proceedings of the 48th International Conference on Parallel Processing (ICPP)*, Kyoto, Japan, August 5-8, 2019.
- [C13] Zhao Lucis Li, Mike Chieh-Jan Liang, **Wei Bai**, Qiming Zheng, Yongqiang Xiong, Guangzhong Sun, "Accelerating Rule-matching Systems with Learned Rankers" in *Proceedings of the 2019 USENIX Annual Technical Conference (ATC)*, Renton, Washington, July 10-12, 2019 (Short Paper).
- [C12] Hong Zhang, Junxue Zhang, **Wei Bai**, Kai Chen, Mosharaf Chowdhury, "Resilient Datacenter Load Balancing in the Wild" in *Proceedings of the ACM SIGCOMM 2017 Conference (SIGCOMM)*, Los Angeles, California, August 21-25, 2017.
- [C11] Ziyang Li, **Wei Bai**, Kai Chen, Dongsu Han, Yiming Zhang, Dongsheng Li, Hongfang Yu, "Rate-Aware Flow Scheduling for Commodity Data Center Networks" in *Proceedings of the 36th Annual IEEE International Conference on Computer Communications (INFOCOM)*, Atlanta, Georgia, May 1-4, 2017.
- [C10] **Wei Bai**, Kai Chen, Li Chen, Changhoon Kim, Haitao Wu, "Enabling ECN over Generic Packet Scheduling" in *Proceedings of the 12th International Conference on emerging Networking EXperiments and Technologies (CoNEXT)*, Irvine, California, December 12-15, 2016.
- [C9] Li Chen, Kai Chen, **Wei Bai**, Mohammad Alizadeh "Scheduling Mix-flows in Commodity Datacenters with Karuna" in *Proceedings of the ACM SIGCOMM 2016 Conference (SIGCOMM)*, Florianopolis, Brazil, August 22-26, 2016.
- [C8] Shuihai Hu, **Wei Bai**, Kai Chen, Chen Tian, Ying Zhang, Haitao Wu, "Providing Bandwidth Guarantees, Work Conservation and Low Latency Simultaneously in the Cloud", in *Proceedings of the 35th Annual IEEE International Conference on Computer Communications (INFOCOM)*, San Francisco, California, April 10-14, 2016.
- [C7] **Wei Bai**, Li Chen, Kai Chen, Haitao Wu "Enabling ECN in Multi-Service Multi-Queue Data Centers" in *Proceedings of the 13th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, Santa Clara, California, March 16-18, 2016.

- [C6] **Wei Bai**, Li Chen, Kai Chen, Dongsu Han, Chen Tian, Hao Wang, "Information-Agnostic Flow Scheduling for Commodity Data Centers" in *Proceedings of the 12th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, Oakland, California, May 4-6, 2015.
- [C5] Shuihai Hu, Kai Chen, Haitao Wu, **Wei Bai**, Chang Lan, Hao Wang, Hongze Zhao, Chuanxiong Guo, "Explicit Path Control in Commodity Data Centers: Design and Applications" in *Proceedings of the 12th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, Oakland, California, May 4-6, 2015.
- [C4] Yangming Zhao, Kai Chen, **Wei Bai**, Minlan Yu, Chen Tian, Yanhui Geng, Yiming Zhang, Dan Li, Sheng Wang, "RAPIER: Integrating Routing and Scheduling for Coflow-aware Data Center Networks" in *Proceedings of the 34th Annual IEEE International Conference on Computer Communications (INFOCOM)*, Hong Kong, April 26-May 1, 2015.
- [C3] Hong Zhang, Kai Chen, **Wei Bai**, Dongsu Han, Chen Tian, Hao Wang, Haibing Guan, Ming Zhang, "Guaranteeing Deadlines for Inter-Datacenter Transfers" in *Proceedings of the 10th European Conference on Computer Systems (EuroSys)*, Bordeaux, France, April 21-24, 2015.
- [C2] **Wei Bai**, Kai Chen, Haitao Wu, Wuwei Lan, Yangming Zhao, "PAC: Taming TCP Incast Congestion Using Proactive ACK Control" in *Proceedings of the IEEE 22nd International Conference on Network Protocols (ICNP)*, Research Triangle, North Carolina, October 21-24, 2014.
- [C1] Yang Peng, Kai Chen, Guohui Wang, **Wei Bai**, Zhiqiang Ma, Lin Gu, "Hadoop-Watch: A First Step Towards Comprehensive Traffic Forecasting in Cloud Computing" in *Proceedings of the 33rd Annual IEEE International Conference on Computer Communications (INFOCOM)*, Toronto, Canada, April 27-May 2, 2014.

Workshop Publications

- [W6] Xinhao Kong, Jiaqi Lou, **Wei Bai**, Nam Sung Kim, Danyang Zhuo, "Towards a Manageable Intra-Host Network" in *Proceedings of the 19th Workshop on Hot Topics in Operating Systems (HotOS)*, Providence, Rhode Island, June 22-24, 2023.
- [W5] Jiacheng Xia, Gaoxiong Zeng, Junxue Zhang, Weiyan Wang, **Wei Bai**, Junchen Jiang, Kai Chen, "Rethinking transport layer design for distributed machine learning" in *Proceedings of the 3rd Asia-Pacific Workshop on Networking (APNet)*, Beijing, China, August 17-18, 2019.
- [W4] Shuihai Hu, **Wei Bai**, Baochen Qiao, Kai Chen, Kun Tan, "Augmenting Proactive Congestion Control with Aeolus" in *Proceedings of the 2nd Asia-Pacific Workshop on Networking (APNet)*, Beijing, China, August 2-3, 2018.
- [W3] Gaoxiong Zeng, **Wei Bai**, Ge Chen, Kai Chen, Dongsu Han, Yibo Zhu, "Combining ECN and RTT for Datacenter Transport" in *Proceedings of the 1st Asia-Pacific Workshop on Networking (APNet)*, Hong Kong, China, August 3-4, 2017.

- [W2] **Wei Bai**, Kai Chen, Shuihai Hu, Kun Tan, Yongqiang Xiong, "Congestion Control for High-speed Extremely Shallow-buffered Datacenter Networks" in *Proceedings of the 1st Asia-Pacific Workshop on Networking (APNet)*, Hong Kong, China, August 3-4, 2017.
- [W1] **Wei Bai**, Li Chen, Kai Chen, Dongsu Han, Chen Tian, Weicheng Sun, "PIAS: Practical Information-Agnostic Flow Scheduling for Data Center Networks" in *Proceedings of the 13th ACM Workshop on Hot Topics in Networks (HotNets)*, Los Angeles, California, October 27-28, 2014.

Journal Publications

- [J10] Gaoxiong Zeng, **Wei Bai**, Ge Chen, Kai Chen, Dongsu Han, Yibo Zhu, Lei Cui, "Congestion Control for Cross-Datacenter Networks", in *IEEE/ACM Transactions on Networking (ToN)*, 2022.
- [J9] Shuihai Hu, Gaoxiong Zeng, **Wei Bai**, Zilong Wang, Baochen Qiao, Kai Chen, Kun Tan, Yi Wang, "Aeolus: A Building Block for Proactive Transport in Datacenter Networks", in *IEEE/ACM Transactions on Networking (ToN)*, 2021.
- [J8] Yang Cheng, Dan Li, Zhiyuan Guo, Binyao Jiang, Jinkun Geng, **Wei Bai**, Jianping Wu, Yongqiang Xiong "Accelerating End-to-End Deep Learning Workflow With Codesign of Data Preprocessing and Scheduling", in *IEEE Transactions on Parallel and Distributed Systems (TPDS)*, 2020.
- [J7] **Wei Bai**, Shuihai Hu, Kai Chen, Kun Tan, Yongqiang Xiong, "One More Config is Enough: Saving (DC)TCP for High-speed Extremely Shallow-buffered Datacenters", in *IEEE/ACM Transactions on Networking (ToN)*, 2020.
- [J6] Shuihai Hu, **Wei Bai**, Kai Chen, Chen Tian, Ying Zhang, Haitao Wu, "Providing Bandwidth Guarantees, Work Conservation and Low Latency Simultaneously in the Cloud", in *IEEE Transactions on Cloud Computing (TCC)*, 2018.
- [J5] Shuhao Liu, Hong Xu, Libin Liu, **Wei Bai**, Kai Chen, Zhiping Cai, "RepNet: Cutting Latency with Flow Replication in Data Center Networks", in *IEEE Transactions on Services Computing (TSC)*, 2018.
- [J4] **Wei Bai**, Li Chen, Kai Chen, Dongsu Han, Chen Tian, Hao Wang, "PIAS: Practical Information-Agnostic Flow Scheduling for Commodity Data Centers" in *IEEE/ACM Transactions on Networking, (ToN)*, 2017.
- [J3] Hong Zhang, Kai Chen, **Wei Bai**, Dongsu Han, Chen Tian, Hao Wang, Haibing Guan, Ming Zhang, "Guaranteeing Deadlines for Inter-Datacenter Transfers" in *IEEE/ACM Transactions on Networking (ToN)*, 2017.
- [J2] Shuihai Hu, Kai Chen, Haitao Wu, **Wei Bai**, Chang Lan, Hao Wang, Hongze Zhao, Chuanxiong Guo, "Explicit Path Control in Commodity Data Centers: Design and Applications" in *IEEE/ACM Transactions on Networking, (ToN)*, 2016.
- [J1] Yang Peng, Kai Chen, Guohui Wang, **Wei Bai**, Yangming Zhao, Hao Wang, Yanhui Geng, Zhiqiang Ma, Lin Gu, "Towards Comprehensive Traffic Forecasting in Cloud Computing: Design and Application" in *IEEE/ACM Transactions on Networking (ToN)*, 2016.

Selected Talks

- Mar. 2024 RDMA at Hyperscale: Experience and Future Directions, King Abdullah University of Science and Technology
- Nov. 2023 RDMA at Hyperscale: Experience and Future Directions, University of Wisconsin-Madison
- Sep. 2023 RDMA at Hyperscale: Experience and Future Directions, Louisiana State University
- June 2023 RDMA at Hyperscale: Experience and Future Directions, APNet 2023, Hong Kong, China
- May 2023 Empowering Azure Storage with RDMA, Peking University
- Apr. 2023 (Guest Lecture) Inside a Datacenter: Network Infrastructure, University of Pittsburgh
- Mar. 2023 Empowering Azure Storage with 100×100 RDMA, University of Washington
- Nov. 2022 (Guest Lecture) Inside a Datacenter: Hardware and Software Architecture, Xiamen University
- Oct. 2022 RDMA Deployments at Microsoft Azure: A Researcher's Experience and View, Intel Labs
- Oct. 2022 (Guest Lecture) Inside a Datacenter: Hardware and Software Architecture, Boston University
- July 2022 Towards Understanding the Micro-Behaviors of Hardware Offloaded Network Stacks, Intel Switch Fabric Group
- June 2022 Towards Understanding the Micro-Behaviors of Hardware Offloaded Network Stacks, Keysight
- May 2022 RDMA at Cloud Region Scale: A Researcher's Experience and View, Systems Seminar, Rice University
- Dec. 2020 Aeolus: A Building Block for Proactive Transport in Datacenters, Systems and Networking Seminar, Duke University
- Jan. 2019 Build Reliable Cloud Networks with SONiC and ONE, OCP China Technology Day, Shenzhen, China
- Oct. 2018 Congestion Control Mechanisms for Data Center Networks, HotDC 2018, Beijing, China
- Aug. 2017 Experiments with Data Center Congestion Control Research, APNet 2017, Hong Kong, China
- Dec. 2016 Enabling ECN over Generic Packet Scheduling, CoNEXT 2016, Irvine, California
- Mar. 2016 Enabling ECN in Multi-Service Multi-Queue Data Centers, NSDI 2016, Santa Clara, California
- May 2015 Information-Agnostic Flow Scheduling for Commodity Data Centers, NSDI 2015, Oakland, California

Professional Activities

Journal Editorship

Editorial Board (Networking Area) Member, Journal of Systems Research (**JSys**), 2021-2022

Technical Program Committee

The 22nd USENIX Symposium on Networked Systems Design and Implementation (**NSDI 2025**)

The 15th ACM SIGOPS Asia-Pacific Workshop on Systems (**APSys 2024**)

The 8th Asia-Pacific Workshop on Networking (**APNet 2024**)

The International Conference for High Performance Computing, Networking, Storage, and Analysis 2024 (**SC 2024**)

The 7th Asia-Pacific Workshop on Networking (**APNet 2023**)

ACM SIGMETRICS 2023

The 6th Asia-Pacific Workshop on Networking (**APNet 2022**)

ACM SIGCOMM Symposium on SDN Research 2022 (**SOSR 2022**)

ACM SIGMETRICS / IFIP Performance 2022 (**SIGMETRICS / Performance 2022**)

ACM SIGCOMM Symposium on SDN Research 2021 (**SOSR 2021**)

The 5th Asia-Pacific Workshop on Networking (**APNet 2021**)

The 4th Asia-Pacific Workshop on Networking (**APNet 2020**)

The 11th ACM SIGOPS Asia-Pacific Workshop on Systems (**APSys 2020**)

The 3rd Asia-Pacific Workshop on Networking (**APNet 2019**)

The 38th IEEE International Conference on Distributed Computing Systems (**ICDCS 2018**), "Cloud Computing and Data Centers" Track

Organizing Committee

Web Chair, ACM SIGCOMM Symposium on SDN Research 2021 (**SOSR 2021**)

Publicity Chair, The 4th Asia-Pacific Workshop on Networking (**APNet 2020**)

Registration Chair, ACM SIGCOMM 2019 Conference (**SIGCOMM 2019**)

Web Chair, The 3rd Asia-Pacific Workshop on Networking (**APNet 2019**)

Web Chair, The 2nd Asia-Pacific Workshop on Networking (**APNet 2018**)

Reviewer

IEEE/ACM Transactions on Networking

ACM Transactions on Computer Systems

Elsevier Computer Networks

IEEE Transactions on Network and Service Management

IEEE Communications Letters

ACM Transactions on Storage

IEEE Transactions on Cloud Computing

ACM Transactions on Architecture and Code Optimization
Elsevier Future Generation Computer Systems



References

Available upon request